

Question1:

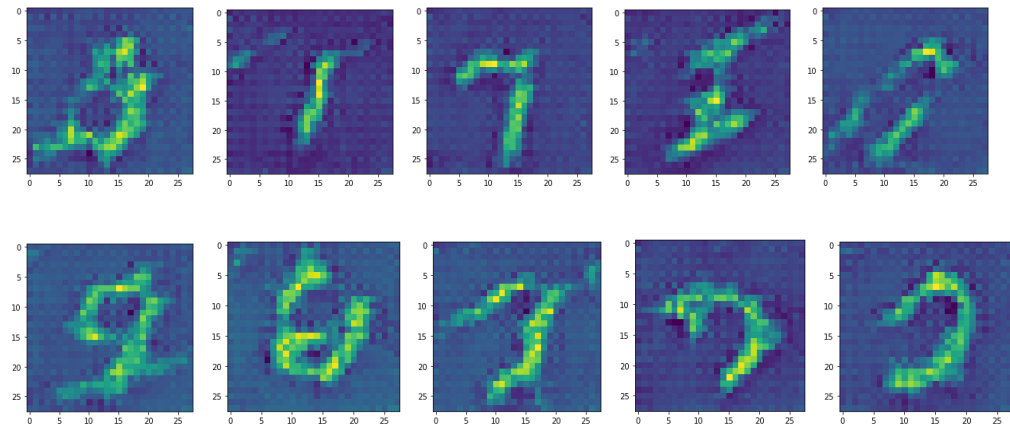
Dataset Used: MNIST

(A) Implementation of DCGAN

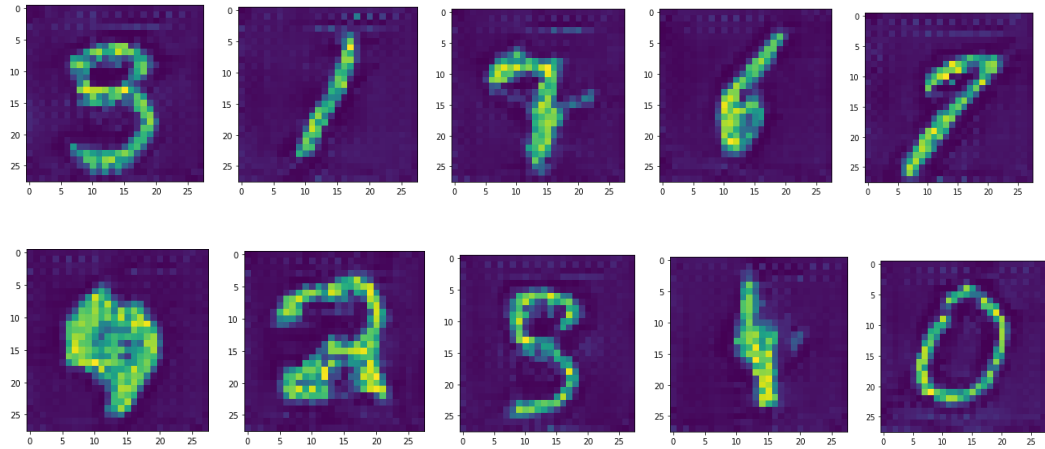
Discriminator Used: Resnet18

Results:

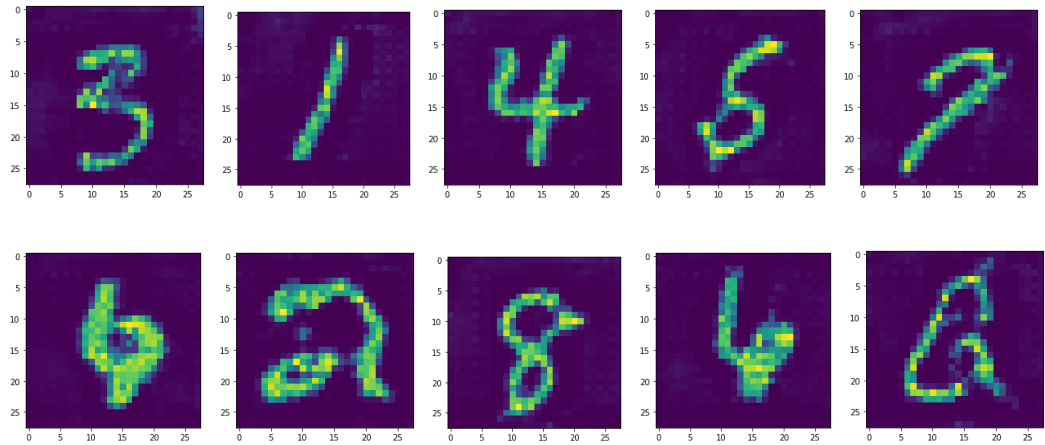
i. Image Generated after 1st epoch



ii. Image Generated after $(n/2)$ 10th epoch



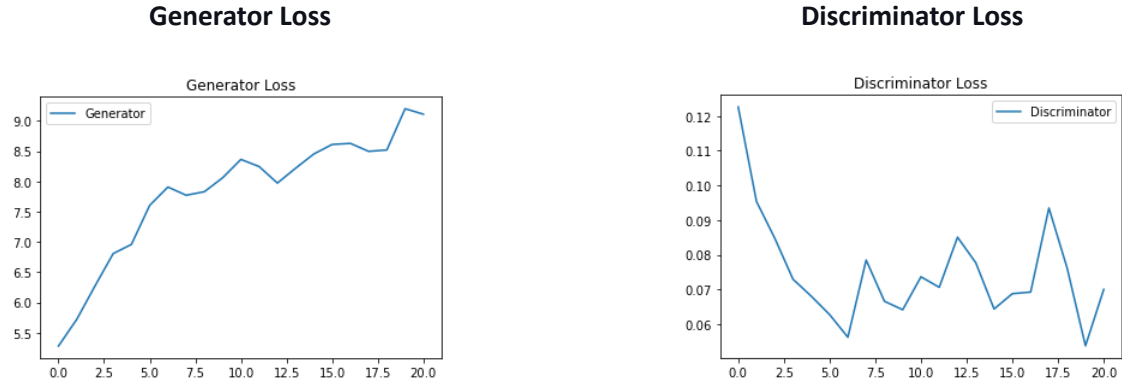
iii. Image Generated after last 20th epoch



Observation:

We can observe that not clear image is generated after 1st epoch the image has many artificats as in 1st epoch Generator has not learnt much. But as we can observe that as epoch continue Generator tries to learn greatly and produce good images and trying to produce images which can fool the discriminator by looking real. After 20th epoch Generator has generated improved quality images some of which may fake the discriminator.

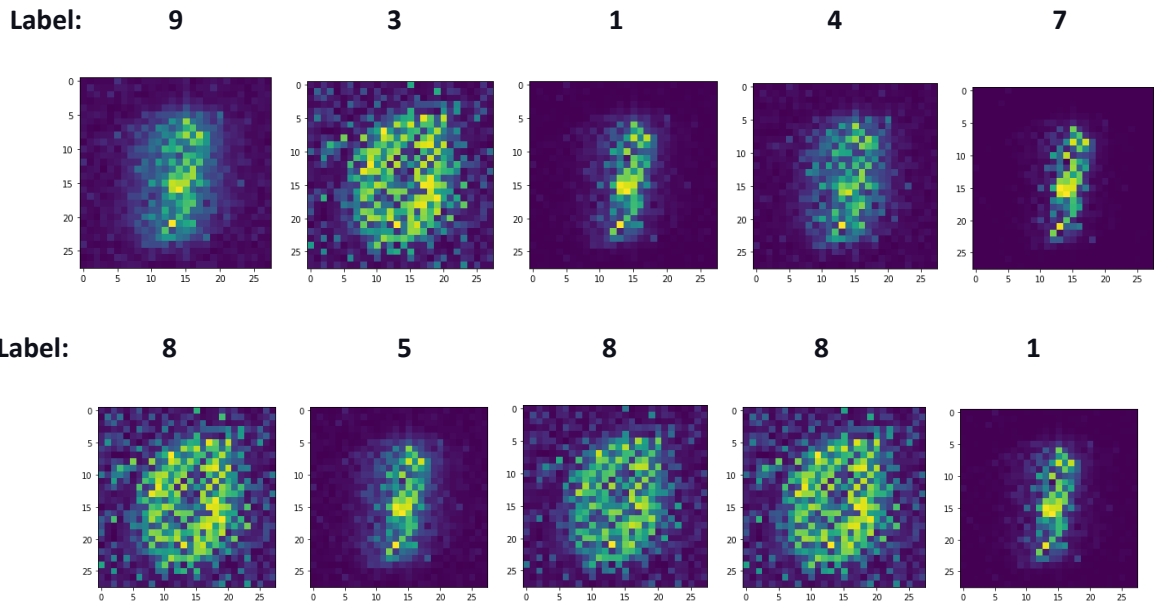
(B) Loss graphs



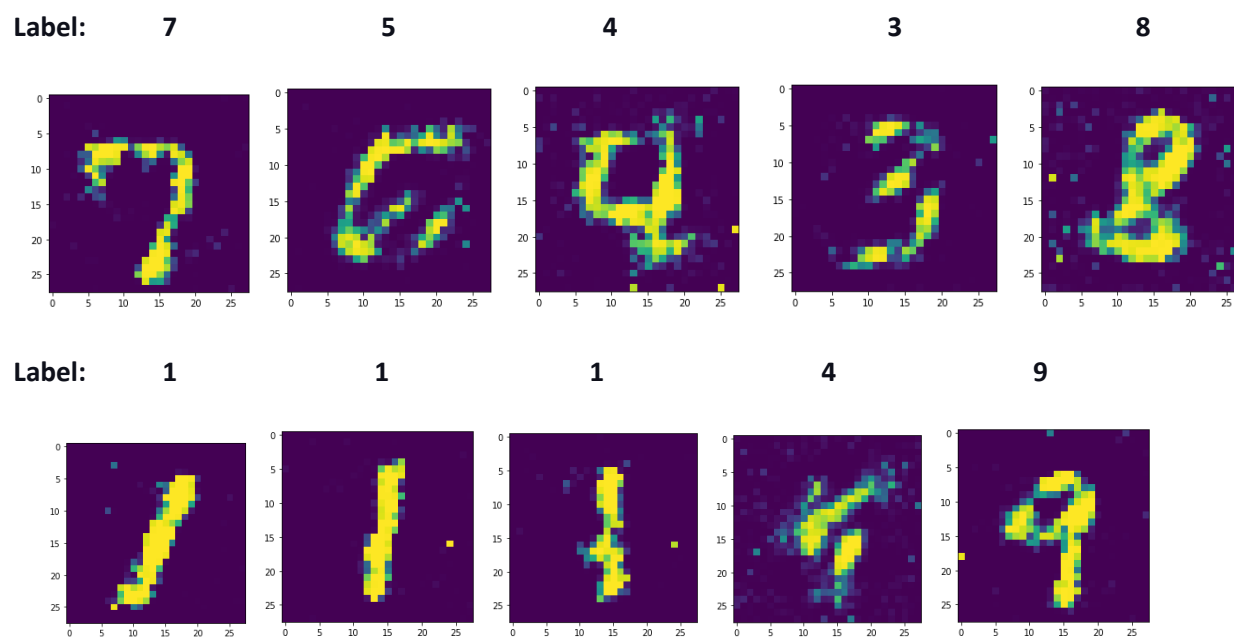
(C) In above model we can't control that which class the generator can produce images. But we can make control on that by using Conditional GANs. The Conditional GANs make the generator to produce class images which it has been asked to by giving particular input as label along with the input vector.

Results:

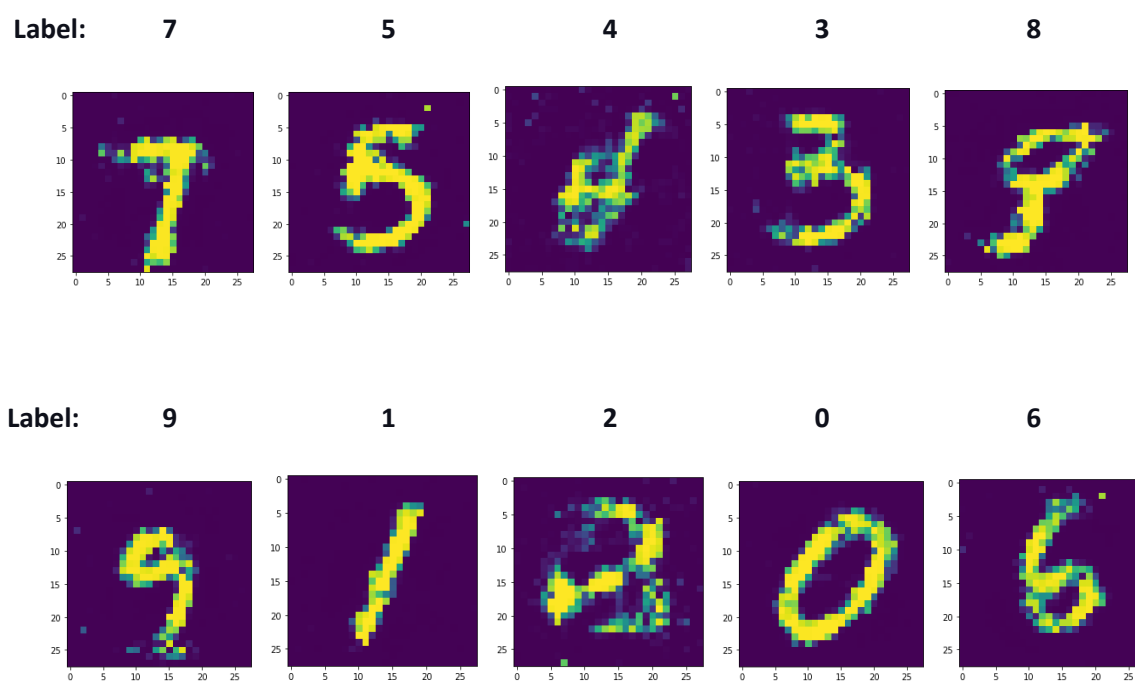
i. Image Generated after 1st epoch



ii. Image Generated after $(n/2)$ 10th epoch



iii. Image Generated after last 20th epoch



Question2:

“Prototypical Networks for Few-shot Learning”

Reproduced results of paper on Dataset: **Omniglot**

Results:

1-shot (5-way)

```
frsd@frsd-DGX-Station: /data/AnnotatedData/single/Disguised...
Avg Train Loss: 0.1148788956925273, Avg Train Acc: 0.9610333341360092
Avg Val Loss: 0.042116488473011485, Avg Val Acc: 0.98480000436306 (Best: 0.99306
66708946227)
=== Epoch: 97 ===
100%|████████████████████████████████████████| 100/100 [00:02<00:00, 47.56it/s]
Avg Train Loss: 0.11764817409217358, Avg Train Acc: 0.9608999955654144
Avg Val Loss: 0.04576633462779, Avg Val Acc: 0.9836000043153763 (Best: 0.9930666
708946227)
=== Epoch: 98 ===
100%|████████████████████████████████████████| 100/100 [00:01<00:00, 53.22it/s]
Avg Train Loss: 0.12416552852839231, Avg Train Acc: 0.9592999976873398
Avg Val Loss: 0.038327422857955755, Avg Val Acc: 0.9884000056982041 (Best: 0.993
0666708946227)
=== Epoch: 99 ===
100%|████████████████████████████████████████| 100/100 [00:01<00:00, 58.09it/s]
Avg Train Loss: 0.12278334762901068, Avg Train Acc: 0.9602000004053116
Avg Val Loss: 0.04525133955503406, Avg Val Acc: 0.9870666694641114 (Best: 0.9930
666708946227)
Testing with last model..
Test Acc: 0.9830933376550675
Testing with best model..
Test Acc: 0.9832933378815651
(GFP) frsd@frsd-DGX-Station: /data/AnnotatedData/single/Disguised GD/pronet/Proto
typical-Networks-for-Few-shot-Learning-PyTorch/src$
```

5-shot (5-way)

```
frsd@frsd-DGX-Station: /data/AnnotatedData/single/Disguised...
frsd@frsd-DGX-Station: /data/AnnotatedData/single/Disguised...
frsd@frsd-DGX-Station: /data/AnnotatedData/single/Disguised...
Avg Train Loss: 0.02871436153771356, Avg Train Acc: 0.9892666751146316
Avg Val Loss: 0.010129360523292785, Avg Val Acc: 0.9964000016450882 (Best: 0.998
2666683197021)
=== Epoch: 97 ===
100%|████████████████████████████████████████| 100/100 [00:02<00:00, 47.89it/s]
Avg Train Loss: 0.029034479069523514, Avg Train Acc: 0.9890000081062317
Avg Val Loss: 0.013357074773009608, Avg Val Acc: 0.9973333346843719 (Best: 0.998
2666683197021)
=== Epoch: 98 ===
100%|████████████████████████████████████████| 100/100 [00:02<00:00, 46.37it/s]
Avg Train Loss: 0.03115166146075353, Avg Train Acc: 0.9890000087022781
Avg Val Loss: 0.011492938051563185, Avg Val Acc: 0.9972000020742416 (Best: 0.998
2666683197021)
=== Epoch: 99 ===
100%|████████████████████████████████████████| 100/100 [00:02<00:00, 47.63it/s]
Avg Train Loss: 0.029068828816525637, Avg Train Acc: 0.989766675233841
Avg Val Loss: 0.008700201934524986, Avg Val Acc: 0.9977333343029022 (Best: 0.998
2666683197021)
Testing with last model..
Test Acc: 0.9949200029969215
Testing with best model..
Test Acc: 0.9954666699767113
(GFP) frsd@frsd-DGX-Station: /data/AnnotatedData/single/Disguised GD/pronet/Proto
typical-Networks-for-Few-shot-Learning-PyTorch/src$
```

1-shot (20-way)

```
frsd@frsd-DGX-Station: /data/AnnotatedData/single/Disguised...
frsd@frsd-DGX-Station: /data/Annota... x frsd@frsd-DGX-Station: /data/Annota... x
Avg Train Loss: 0.12711915116757155, Avg Train Acc: 0.9590333324670791
Avg Val Loss: 0.14776318622753024, Avg Val Acc: 0.9536333352327346 (Best: 0.9625
666707754135)
=== Epoch: 97 ===
100%|████████████████████████████████████████████████████████████████████████████████| 100/100 [00:02<00:00, 42.71it/s]
Avg Train Loss: 0.12194435965269804, Avg Train Acc: 0.959599991893769
Avg Val Loss: 0.13213673869147896, Avg Val Acc: 0.9577999997138977 (Best: 0.9625
666707754135)
=== Epoch: 98 ===
100%|████████████████████████████████████████████████████████████████████████████████| 100/100 [00:02<00:00, 36.67it/s]
Avg Train Loss: 0.1244412138313055, Avg Train Acc: 0.95803333292484283
Avg Val Loss: 0.15098728029057384, Avg Val Acc: 0.9503333353996277 (Best: 0.9625
666707754135)
=== Epoch: 99 ===
100%|████████████████████████████████████████████████████████████████████████████████| 100/100 [00:02<00:00, 42.22it/s]
Avg Train Loss: 0.11588014565408229, Avg Train Acc: 0.9620000004768372
Avg Val Loss: 0.14742135919630528, Avg Val Acc: 0.9508666694164276 (Best: 0.9625
666707754135)
Testing with last model..
Test Acc: 0.9449733336567878
Testing with best model..
Test Acc: 0.9461000009179116
(GFP) frsd@frsd-DGX-Station: /data/AnnotatedData/single/Disguised GD/pronet/Proto
typical-Networks-for-Few-shot-Learning-PyTorch/src$
```

5-shot (20-way)

```
frsd@frsd-DGX-Station: /data/AnnotatedData/single/Disguised...
=== Epoch: 97 ===
100%|████████████████████████████████████████████████████████████████████████████████| 100/100 [00:02<00:00, 43.30it/s]
Avg Train Loss: 0.0281412196950987, Avg Train Acc: 0.989400006532669
Avg Val Loss: 0.03528093245695345, Avg Val Acc: 0.9875000071525574 (Best: 0.9893
00007224083)
=== Epoch: 98 ===
100%|████████████████████████████████████████████████████████████████████████████████| 100/100 [00:02<00:00, 38.73it/s]
Avg Train Loss: 0.028378184132743626, Avg Train Acc: 0.9887333416938782
Avg Val Loss: 0.04044244340155274, Avg Val Acc: 0.9864333415031433 (Best: 0.9893
00007224083)
=== Epoch: 99 ===
100%|████████████████████████████████████████████████████████████████████████████████| 100/100 [00:02<00:00, 39.42it/s]
Avg Train Loss: 0.030092107636155562, Avg Train Acc: 0.988466677069664
Avg Val Loss: 0.0390228965782444, Avg Val Acc: 0.9878333407640457 (Best: 0.98930
0007224083)
Testing with last model..
Test Acc: 0.9842233434319496
Testing with best model..
Test Acc: 0.9851066771745681
PyThreadState_Clear: warning: thread still has a frame
free(): invalid next size (fast)
Aborted (core dumped)
(GFP) frsd@frsd-DGX-Station: /data/AnnotatedData/single/Disguised GD/pronet/Proto
typical-Networks-for-Few-shot-Learning-PyTorch/src$
```

Model	1-shot (5-way Acc.)	5-shot (5-way Acc.)	1-shot (20-way Acc.)	5-shot (20-way Acc.)
Values Reported in paper	98.8%	99.7%	96.0%	98.9%
Values after reproducing paper	98.32%	99.54%	94.61%	98.51%

Analysis:

Above table provides reproduced results on the given dataset “Omniglot”. The model is trained for 100 epochs, which gives an accuracy similar to the results reported in the paper for different types N-way K-shot classification. Also, best model is returned after training for all epochs. Authors uses prototypical network to classify Omniglot data samples by calculating distance between prototype representations with others in order to exhibit Inductive-bias technique.

Explanation of Algorithm:

As we know few shot classification is type of learning technique whenever there is less training data available. Few shot learning is a kind of meta-learning in which the model is trained on several related tasks, during the training phase. It is trained in such a way that it can perform well on non-seen classes.

This algorithm represents that there exist an embedding in which points cluster around a single prototype representation for every class. It has k-way n-shot type of training. Here we have k classes and for each class we have 2 set : support set and query set.

Algorithm steps:

1. Flatten the images to be transformed into 1-D vectors.
2. Then the class prototypes are computed by forming the clusters and each cluster is represented by a centroid.
3. The embeddings of the support set images are averaged to form a class prototype.
4. Now distance is computed between the queries and prototypes. In this the metric choice is very important, the authors have specified their own choice of distance metric.
5. After computing the distance softmax is performed over distances to the prototypes in order to get probabilities.

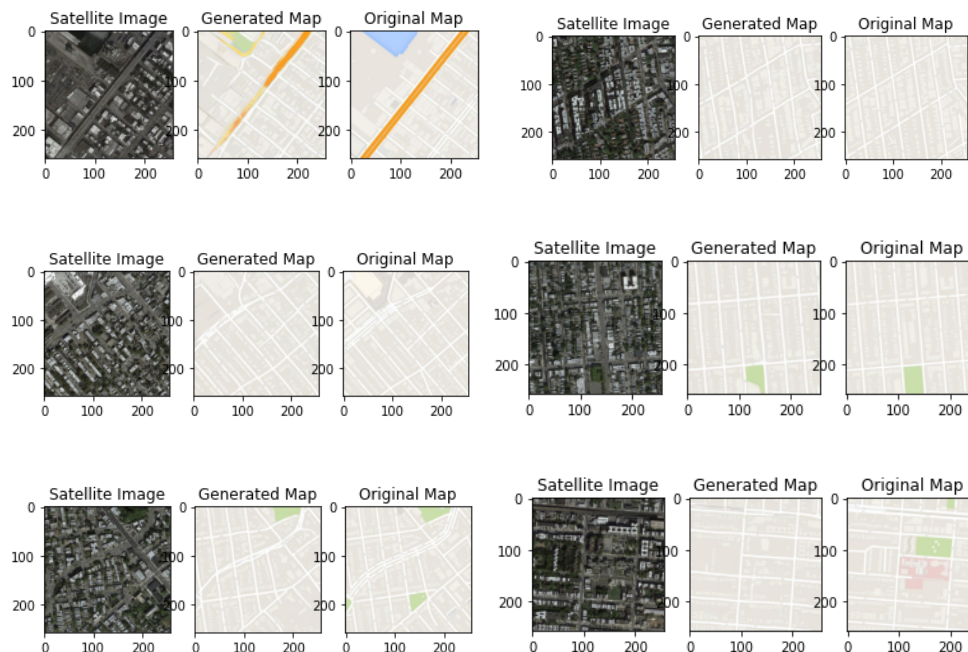
6. Now the classes are classified on the basis of probability scores. The higher the score more the chance of the query sample belonging to that class.
7. Finally the loss is calculated and backpropagated.

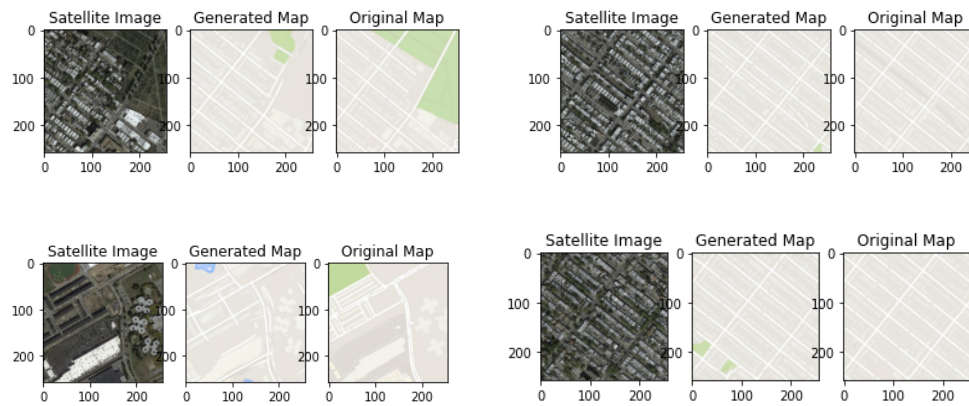
Pseudo code for Reproducing the results of this paper

- The embedding part takes a (28x28x3) image and returns a column vector of length 64.
- The image2vector function is composed of 4 modules.
- Each module consists of various layers: conv layer, batch norm layer, ReLU activation function and 2x2 max pooling layer.
- **Optimizer:** Adam
- **learning rate:** 0.0001
- The model is trained for 2000 episodes for every 5th epoch. training is performed by randomly picking new sample in the training set at each episode.
- It is tested on 1000 episodes.

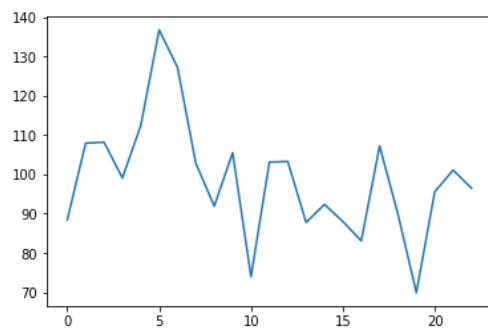
Question3:

(a) Trained GAN generated results are as follows

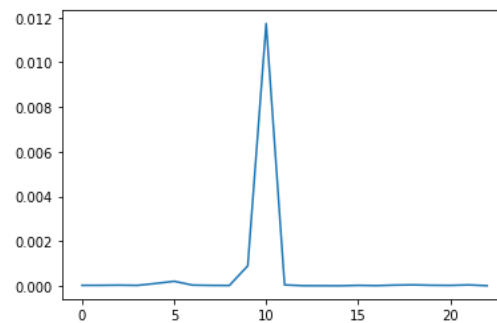




(b) Generator loss graph



Discriminator loss graph



(c) Pretrained GAN and Trained GAN results

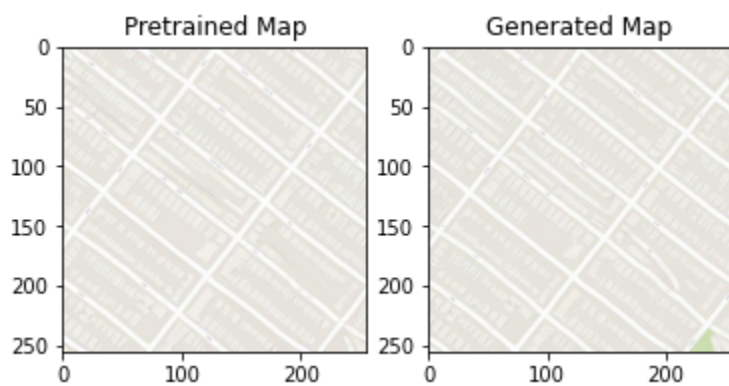
SSIM = 0.686



SSIM = 0.701



SSIM = 0.657



SSIM = 0.688



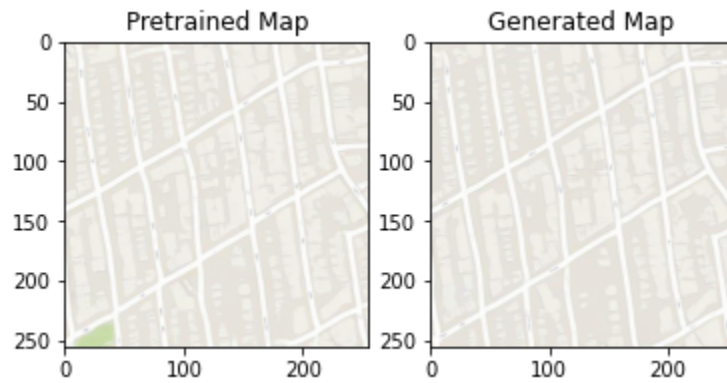
SSIM = 0.668



SSIM = 0.718



SSIM = 0.683



Observations:

Pretrained model used is CycleGAN As we can see from the above SSIM (Structural Similarity Index Measure) pretrained models generate better visual maps of satellite images. Images generated by pretrained GAN are very similar to the ground truth maps as compared to images generated by trained model.

Question4:

“DARTS: Differentiable Architecture Search”

Reproduced results of paper on Dataset: **CIFAR-10**

```
frsd@frsd-DGX-Station: /data/AnnotatedData/single/Disguised...
05/03 12:10:26 PM train_acc 96.523998
05/03 12:10:26 PM valid_000 2.226235e-01 95.833328 100.000000
05/03 12:10:28 PM valid_050 1.364237e-01 96.384801 99.877450
05/03 12:10:30 PM valid_100 1.290726e-01 96.379948 99.927805
05/03 12:10:30 PM valid_acc 96.379997
05/03 12:10:30 PM epoch 99 lr 0.000000e+00
05/03 12:10:31 PM train_000 2.308999e-01 94.791664 100.000000
05/03 12:10:42 PM train_050 1.745605e-01 96.548200 99.918300
05/03 12:10:54 PM train_100 1.731540e-01 96.503710 99.917491
05/03 12:11:05 PM train_150 1.706209e-01 96.523176 99.931015
05/03 12:11:17 PM train_200 1.685042e-01 96.641788 99.937811
05/03 12:11:29 PM train_250 1.721872e-01 96.518092 99.941899
05/03 12:11:40 PM train_300 1.723405e-01 96.497783 99.951550
05/03 12:11:52 PM train_350 1.729767e-01 96.504033 99.958452
05/03 12:12:03 PM train_400 1.730547e-01 96.464565 99.958437
05/03 12:12:15 PM train_450 1.720428e-01 96.484661 99.960735
05/03 12:12:27 PM train_500 1.734915e-01 96.463321 99.962575
05/03 12:12:31 PM train_acc 96.459997
05/03 12:12:31 PM valid_000 2.132604e-01 94.791664 100.000000
05/03 12:12:33 PM valid_050 1.364635e-01 96.200978 99.918300
05/03 12:12:36 PM valid_100 1.303115e-01 96.225245 99.948432
05/03 12:12:36 PM valid_acc 96.219998
(GFP) frsd@frsd-DGX-Station: /data/AnnotatedData/single/Disguised GD/pronet/darts
/cnn$
```

DARTS: Differentiable Architecture Search

- This paper takes a different approach to the problem and propose DARTS, an efficient architecture search method" (Differentiable Architecture Search). Instead of searching over a discrete set of candidate architectures, we broaden the search space to be continuous, allowing the architecture to be optimised by gradient descent with respect to its validation set performance.
- DARTS achieves competitive performance with the state of the art using orders of magnitude fewer computation resources due to the data efficiency of gradient-based optimization as opposed to inefficient black-box search.
- Authors present a novel bilevel optimization-based algorithm for differentiable network architecture search.

Algorithm 1: DARTS – Differentiable Architecture Search

Create a mixed operation $\bar{o}^{(i,j)}$ parametrized by $\alpha^{(i,j)}$ for each edge (i, j)

while not converged do

1. Update architecture α by descending $\nabla_{\alpha} \mathcal{L}_{val}(w - \xi \nabla_w \mathcal{L}_{train}(w, \alpha), \alpha)$
($\xi = 0$ if using first-order approximation)
2. Update weights w by descending $\nabla_w \mathcal{L}_{train}(w, \alpha)$

Derive the final architecture based on the learned α .

Explanation for Above Algorithm:

1. The problem with searching over a discrete set of candidate operations is that the model has to be trained on specific configuration before moving onto the next configuration. This obviously is time-consuming. Authors found a way of relaxing the discrete set of candidate operations.
2. This is the model's core structure. There are one or more nodes in the cell. These nodes are also referred to as states.
3. The input to a cell is the output of the previous two cells, just like ResNets. There are nodes in this cell. Assume we create a cell with three states/nodes. As a result, the first node will have two inputs, i.e. outputs from the previous two cells.
4. The second state will receive inputs from the first state as well as outputs from the last two cells, for a total of three inputs.

5. The third state will receive inputs from the second and first states, as well as outputs from the last.
6. They used Bilevel Optimization

$$\begin{aligned} \min_{\alpha} \quad & \mathcal{L}_{\text{val}}(w^*(\alpha), \alpha) \\ \text{s.t.} \quad & w^*(\alpha) = \operatorname{argmin}_w \mathcal{L}_{\text{train}}(w, \alpha) \end{aligned}$$

Since we have weight available that have been optimised on the training set, the optimisation problem can be framed as finding the alphas to minimise validation loss.

$$\nabla_{\alpha} L_{\text{val}}(w^*(\alpha), \alpha)$$

α : parameter for operation weight $w^*(\alpha)$: optimal conv weights for specific α

7. Obtaining the optimal weight for each configuration of alpha in equation above requires two optimization loops, hence the authors suggested approximating in such a way that does not need to be optimised until convergence.

Steps to run code:

- For training

```
%cd cnn
```

```
!python train.py --auxiliary --cutout
```

- For Testing

```
%cd cnn
```

```
!python test.py --auxiliary --model_path weights.pt
```

****Note:** Question2 and Question 3 have run the github so not providing the code in the submission.

